

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from matplotlib.dates import DateFormatter
import statsmodels.api as sm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
In [2]: def read_data(path, date="Date"):
    data = pd.read_csv(path)
    data[date] = pd.to_datetime(data[date], dayfirst=True)
    return data
```

```
In [3]: #import csv data
nvda = read_data("data/NVDA.csv")
#nvda['Close'] = nvda['Close'].apply(np.log)
display(nvda)
```

	Date	Open	High	Low	Close	Adj Close	Volume
0	2019-04-15	47.424999	47.615002	45.775002	46.575001	46.258690	156704800
1	2019-04-22	46.337502	48.202499	43.325001	44.522499	44.220123	231328000
2	2019-04-29	44.602501	46.217499	43.875000	45.752499	45.441769	184062800
3	2019-05-06	43.875000	45.084999	41.000000	42.205002	41.918369	262631200
4	2019-05-13	40.875000	41.107498	38.735001	39.132500	38.866734	348726800
...
257	2024-03-18	903.880005	947.780029	850.099976	942.890015	942.890015	288579700
258	2024-03-25	939.409973	967.659973	891.229980	903.559998	903.559998	208706300
259	2024-04-01	902.989990	922.250000	858.799988	880.080017	880.080017	208939400
260	2024-04-08	887.000000	907.390015	830.219971	881.859985	881.859985	207522200
261	2024-04-15	890.979980	906.130005	859.289978	860.010010	860.010010	44307700

262 rows × 7 columns

```
In [4]: searchvol = read_data("data/ai-timeline_Glimpse_Google-Trends.csv", date='T
lnsearchvol = read_data("data/ai-timeline_Glimpse_Google-Trends.csv", date=

lnsearchvol['Absolute Google Search Volume'] = lnsearchvol['Absolute Google

display(searchvol)
display(lnsearchvol)
```

	Time (week of)	Normalized Value (0-100)	Absolute Google Search Volume
0	2019-04-14	15	1360594
1	2019-04-21	14	1269888
2	2019-04-28	15	1360594
3	2019-05-05	15	1360594
4	2019-05-12	15	1360594
...
257	2024-03-17	98	8889219
258	2024-03-24	100	9070631
259	2024-03-31	96	8707806
260	2024-04-07	97	8798512
261	2024-04-14	97	8798512

262 rows × 3 columns

	Time (week of)	Normalized Value (0-100)	Absolute Google Search Volume
0	2019-04-14	15	14.123432
1	2019-04-21	14	14.054439
2	2019-04-28	15	14.123432
3	2019-05-05	15	14.123432
4	2019-05-12	15	14.123432
...
257	2024-03-17	98	16.000350
258	2024-03-24	100	16.020552
259	2024-03-31	96	15.979730
260	2024-04-07	97	15.990093
261	2024-04-14	97	15.990093

262 rows × 3 columns

In [5]:

```
#cleaning data
nvda = nvda.drop(columns=['Open', 'High', 'Low', 'Adj Close'])
display(nvda)
```

	Date	Close	Volume
0	2019-04-15	46.575001	156704800
1	2019-04-22	44.522499	231328000
2	2019-04-29	45.752499	184062800
3	2019-05-06	42.205002	262631200
4	2019-05-13	39.132500	348726800
...
257	2024-03-18	942.890015	288579700
258	2024-03-25	903.559998	208706300
259	2024-04-01	880.080017	208939400
260	2024-04-08	881.859985	207522200
261	2024-04-15	860.010010	44307700

262 rows × 3 columns

```
In [6]: searchvol = searchvol.rename(columns = {'Time (week of)': 'Date',"Absolute
Insearchvol = Insearchvol.rename(columns = {'Time (week of)': 'Date',"Absol
display(searchvol)
display(Insearchvol)
```

	Date	Normalized Value (0-100)	Search Volume
0	2019-04-14	15	1360594
1	2019-04-21	14	1269888
2	2019-04-28	15	1360594
3	2019-05-05	15	1360594
4	2019-05-12	15	1360594
...
257	2024-03-17	98	8889219
258	2024-03-24	100	9070631
259	2024-03-31	96	8707806
260	2024-04-07	97	8798512
261	2024-04-14	97	8798512


262 rows × 3 columns

	Date	Normalized Value (0-100)	Search Volume
0	2019-04-14	15	14.123432
1	2019-04-21	14	14.054439
2	2019-04-28	15	14.123432
3	2019-05-05	15	14.123432
4	2019-05-12	15	14.123432
...
257	2024-03-17	98	16.000350
258	2024-03-24	100	16.020552
259	2024-03-31	96	15.979730
260	2024-04-07	97	15.990093
261	2024-04-14	97	15.990093


262 rows × 3 columns

```
In [7]: def plot_uni(data,variable,colour):  
plt.figure(figsize=(14, 8))  
plt.xlabel('Date')  
plt.ylabel(variable)  
plt.title('Univariate '+variable+' Analysis')  
# Format x-axis dates  
date_form = DateFormatter("%d/%m/%Y")  
plt.gca().xaxis.set_major_formatter(date_form)  
plt.plot(data['Date'], data[variable], color=colour)
```

```
In [8]: plot_uni(nvda,'Close','blue')  
plt.savefig("outputs/figure1.png")  
plt.show()
```



```
In [9]: plot_uni(searchvol,'Search Volume','red')  
plt.savefig("outputs/figure2")  
plt.show()
```



```
In [10]: def scientific_formatter(x, pos):
# Format the tick label in scientific notation with exponent next to the
return "{:.0e}".format(x)


In [11]: def plot_comparison(nvda, searchvol, important_dates=None):
# Plotting historical close prices and search volume
plt.figure(figsize=(14, 8))
plt.xlabel('Date')
plt.ylabel('Close Price')
plt.title('Historical Close Prices and Search Volume')
# Format x-axis dates
date_form = DateFormatter("%d/%m/%Y")
plt.gca().xaxis.set_major_formatter(date_form)

# Plot Close Price on primary y-axis
stock = plt.plot(nvda['Date'], nvda['Close'], label='Close Price', color='blue')
# Create secondary y-axis for Search Volume
ax2 = plt.twinx()
search = ax2.plot(searchvol['Date'], searchvol['Search Volume'], label='Search Volume', color='red')
ax2.set_ylabel('Search Volume')


# Combine legend for both primary and secondary plots
lns = stock + search
labs = [l.get_label() for l in lns]
plt.legend(lns, labs, loc='upper left')

if important_dates != None:
    for item in important_dates:
        plt.axvline(x=pd.Timestamp(item['date']), color='gray', linestyle='dashed')
        y = plt.ylim()[0] + (.5*plt.ylim()[1]-6.75) if item['top'] == 'T' else None
        plt.text(pd.Timestamp(item['date']), y, item['name'], rotation=90)

In [12]: plot_comparison(nvda, searchvol)
plt.savefig("outputs/figure3")
plt.show()
```




```
In [13]: plot_comparison(nvda, lnsearchvol)
plt.show()
```




```
In [14]: # Dates of major AI releases
important_dates = [
    {'date': '2024-02-15', 'name': 'Sora', 'top': False},
    {'date': '2023-12-6', 'name': 'Google Gemini', 'top': False},
    {'date': '2023-11-1', 'name': 'Microsoft Co-pilot', 'top': False},
    {'date': '2022-11-30', 'name': 'ChatGPT', 'top': True},
    {'date': '2021-10-27', 'name': 'GitHub Co-pilot', 'top': True}
] # Example dates and corresponding titles

plot_comparison(nvda, searchvol, important_dates)
plt.savefig("outputs/figure4")
plt.show()
```



```
In [15]: plot_comparison(nvda, lnsearchvol, important_dates)
plt.savefig("outputs/figure7")
plt.show()
```



```
In [16]: # x_train, x_test, y_train, y_test = train_test_split(joined['Search Volume
# plt.scatter(x_train, y_train, label='Trainng dtaa',color='r',alpha=.7)
# plt.scatter(x_test, y_test, label='Testing dtaa',color='b',alpha=.4)
# plt.xlabel('search volume')
# plt.ylabel('stock close price')
# plt.legend()
# plt.title('test train split')
# plt.show()

# LR = LinearRegression()
# LR.fit(x.values.reshape(-1,1),y)
# predict = LR.predict(x.values.reshape(-1,1))

# plt.plot(x,predict,label='linear',color='b')
# plt.scatter(x,y,label='actual test data',color='g',alpha=.7)
# plt.xlabel('search volume')
# plt.ylabel('stock close price')
# plt.legend()
# plt.show()
# b= LR.coef_
# print(b)
# #R-squared of model
# LR.score(x.values.reshape(-1,1),y.values)
```




```
In [17]: def calc_regression(x, y):  
    x_resaped = x.values.reshape(-1, 1)  
  
    # Create and fit the linear regression model  
    LR = LinearRegression()  
    LR.fit(x_resaped, y)  
  
    # Predict using the fitted model  
    predict = LR.predict(x_resaped)  
    # Display the coefficient (slope) of the linear regression model  
    coef = LR.coef_  
    print("Coefficient (slope) of the linear regression model:", coef)  
  
    # Calculate R-squared (coefficient of determination) of the model  
    r_squared = LR.score(x_resaped, y)  
    print("R-squared of the linear regression model:", r_squared)  
  
    return predict
```

```
In [18]: def plot_regression(x, y, predict):  
    # Plotting the linear regression line and actual data  
    plt.plot(x, predict, label='Linear Regression', color='r')  
    plt.scatter(x, y, label='Actual Data', color='dimgray', alpha=1, linewidth=1)  
    plt.xlabel('Search Volume')  
    plt.ylabel('Stock Close Price')  
    plt.legend()  
    plt.title('Linear Regression: Search Volume vs. Stock Close Price')
```


```
In [19]: predict1 = calc_regression(searchvol['Search Volume'], nvda['Close'])  
plot_regression(searchvol['Search Volume'], nvda['Close'], predict1)  
plt.savefig("outputs/figure5")  
plt.show()
```

Coefficient (slope) of the linear regression model: $[6.57219112e-05]$
R-squared of the linear regression model: 0.7786908923901182



```
In [20]: predict2 = calc_regression(lnsearchvol['Search Volume'], nvda['Close'])
plot_regression(lnsearchvol['Search Volume'], nvda['Close'], predict2)
plt.savefig("outputs/figure8")
plt.show()
```

Coefficient (slope) of the linear regression model: [234.11895979]
R-squared of the linear regression model: 0.7247326809795085



```
In [21]: def save_regression(figure):
# Get the summary as a string
summary_str = results.summary().as_text()
# Define the file path where you want to save the summary
output_file_path = "regression_outputs/"+figure+".txt"
# Write the summary string to a text file
with open(output_file_path, 'w') as f:
    f.write(summary_str)
```

```
In [22]: results = sm.OLS(nvda['Close'].values, sm.add_constant(searchvol['Search Vo
print(results.summary())
save_regression("figure6")
```

OLS Regression Results						
=====						
=====						
Dep. Variable:	y	R-squared:				
0.779						
Model:	OLS	Adj. R-squared:				
0.778						
Method:	Least Squares	F-statistic:	9			
14.8						
Date:	Wed, 24 Apr 2024	Prob (F-statistic):	3.96			
e-87						
Time:	21:46:22	Log-Likelihood:	-15			
43.0						
No. Observations:	262	AIC:	3			
090.						
Df Residuals:	260	BIC:	3			
097.						
Df Model:	1					
Covariance Type:	nonrobust					
=====						
=====						
	coef	std err	t	P> t	[0.025	0.
975]						

const	28.4688	8.366	3.403	0.001	11.995	4
4.943						
x1	6.572e-05	2.17e-06	30.246	0.000	6.14e-05	7
e-05						
=====						
=====						
Omnibus:	40.594	Durbin-Watson:				
0.064						
Prob(Omnibus):	0.000	Jarque-Bera (JB):	6			
3.567						
Skew:	0.903	Prob(JB):	1.57			
e-14						
Kurtosis:	4.600	Cond. No.	5.94			
e+06						
=====						
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.94e+06. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [23]: results = sm.OLS(nvda['Close'].values, sm.add_constant(lnsearchvol['Search
print(results.summary())
save_regression("figure9")
```

OLS Regression Results

```
=====
====
Dep. Variable:          y      R-squared:
0.725
Model:                OLS      Adj. R-squared:
0.724
Method:                Least Squares      F-statistic:          6
84.5
Date:                  Wed, 24 Apr 2024      Prob (F-statistic):      8.54
e-75
Time:                  21:46:22      Log-Likelihood:          -15
71.6
No. Observations:      262      AIC:          3
147.
Df Residuals:          260      BIC:          3
154.
Df Model:              1
Covariance Type:       nonrobust
=====
====
              coef      std err          t      P>|t|      [0.025      0.
975]
-----
----
const      -3201.6927      130.967      -24.447      0.000      -3459.584      -294
3.801
x1           234.1190         8.948       26.164      0.000       216.499       25
1.739
=====
====
Omnibus:              66.055      Durbin-Watson:
0.052
Prob(Omnibus):         0.000      Jarque-Bera (JB):          15
3.230
Skew:                  1.194      Prob(JB):          5.33
e-34
Kurtosis:              5.887      Cond. No.
319.
=====
====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
In [24]: SnP = read_data("data/SPY.csv")
#cleaning data
SnP = SnP.drop(columns=['Open', 'High', 'Low', 'Adj Close'])
display(SnP)
```

	Date	Close	Volume
0	2019-04-15	290.019989	228726700
1	2019-04-22	293.410004	251486900
2	2019-04-29	294.029999	331555200
3	2019-05-06	288.100006	559396700
4	2019-05-13	285.839996	455352700
...
257	2024-03-18	521.210022	358522300
258	2024-03-25	523.070007	293270500
259	2024-04-01	518.429993	367203800
260	2024-04-08	510.850006	361747100
261	2024-04-15	504.450012	92101400

262 rows × 3 columns

```
In [25]: multi= pd.concat([lnsearchvol['Search Volume'],SnP['Close']],axis=1)
```

```
In [26]: predict3 = calc_regression(SnP['Close'], nvda['Close'])
plot_regression(SnP['Close'], nvda['Close'], predict3)
```

Coefficient (slope) of the linear regression model: [2.31572742]
R-squared of the linear regression model: 0.651732842120399



```
In [27]: multi_resaped = multi.values
multi_resaped
# Create and fit the linear regression model
LR = LinearRegression()
LR.fit(multi_resaped, nvda['Close'])


# Predict nvidia price based on ai search volume and s&p 500 regression
predict4 = LR.predict(multi_resaped)
```

In [28]:

```
# Plotting the Linear regression lines adn
plt.plot(np.log(predict4/predict4[0]), label='MLR', color='r')
plt.plot(np.log(SnP['Close']/SnP['Close'][0]),label='S&P 500 ',color='b')
plt.plot(np.log(nvda['Close']/nvda['Close'][0]),label='nvidia',color='g')
#plt.xlabel('time')
plt.ylabel('Stock Close Price index (base 0)')
plt.legend()
plt.title('Index Logarithmic Price Comparison ')
plt.savefig("outputs/figure11")
plt.show()
```

C:\Users\jarvit\AppData\Local\Temp\ipykernel_20272\2289097429.py:2: RuntimeWarning: invalid value encountered in log

```
plt.plot(np.log(predict4/predict4[0]), label='MLR', color='r')
```




```
In [29]: results= sm.OLS(nvda['Close'].values, sm.add_constant(multi).values).fit()
print(results.summary())
save_regression("figure10")
```

OLS Regression Results

```
=====
====
Dep. Variable:          y      R-squared:
0.848
Model:                OLS      Adj. R-squared:
0.846
Method:              Least Squares      F-statistic:          7
20.5
Date:                Wed, 24 Apr 2024      Prob (F-statistic):      1.50e
-106
Time:                21:46:22      Log-Likelihood:          -14
94.1
No. Observations:      262      AIC:                2
994.
Df Residuals:          259      BIC:                3
005.
Df Model:              2
Covariance Type:      nonrobust
=====
====
              coef      std err          t      P>|t|      [0.025      0.
975]
-----
----
const      -2566.9480      107.041      -23.981      0.000      -2777.730      -235
6.166
x1           156.3712       8.568       18.250      0.000       139.499       17
3.243
x2           1.2919       0.089       14.456      0.000        1.116
1.468
=====
====
Omnibus:              103.416      Durbin-Watson:
0.057
Prob(Omnibus):        0.000      Jarque-Bera (JB):        44
6.523
Skew:                 1.590      Prob(JB):                1.09
e-97
Kurtosis:             8.549      Cond. No.                9.39
e+03
=====
====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 9.39e+03. This might indicate that there are strong multicollinearity or other numerical problems.

In []:

